



— BioMistral-7B
 — OpenBioLLM-8B
 — Mistral-7B-Instruct
 — Base + UQ
 - - NS-AI + UQ
 ⋯⋯ NS-AI + UQ + CL



Deployed τ (markers). Each τ is the largest threshold whose one-sided Clopper-Pearson upper bound ($\delta = 0.10$) on the calibration split's selective error is $\leq \alpha = 0.10$. The bound is evaluated point-wise, once per candidate τ , during selection; it is not a band on the curves above, which are test and held-out data. — : the gate decides every held-out item.

Model	Row	calibration at τ			test at τ		held-out at τ	
		n	risk	CP upper none $\leq \alpha$	defers	risk	defers	risk
BioMistral-7B	Base + UQ	66,938	—	—	100.0%	($\tau=\infty$)	100.0%	($\tau=\infty$)
BioMistral-7B	NS-AI + UQ	41,140	0.024	0.030	96.5%	0.019	—	—
BioMistral-7B	NS-AI + UQ + CL	41,140	0.092	0.094	51.9%	0.090	—	—
OpenBioLLM-8B	Base + UQ	66,938	0.062	0.066	91.9%	0.062	100.0%	(τ < all u)
OpenBioLLM-8B	NS-AI + UQ	3,827	0.000	0.006	99.2%	0.000	—	—
OpenBioLLM-8B	NS-AI + UQ + CL	41,140	0.086	0.092	91.4%	0.101	—	—
Mistral-7B-Instruct	Base + UQ	66,938	0.094	0.099	90.3%	0.091	99.7%	0.000
Mistral-7B-Instruct	NS-AI + UQ	41,140	0.086	0.090	70.2%	0.088	—	—
Mistral-7B-Instruct	NS-AI + UQ + CL	41,140	0.095	0.100	86.2%	0.087	—	—